

---

# Spam Detection

---

## Objective

The objective of this project is to build a machine learning model that can classify text messages as either spam or ham (non-spam). The model uses the Multinomial Naive Bayes algorithm, which is effective for text classification tasks.

## Tools:

This project is built using the following Python libraries:

- **pandas**: For data manipulation and analysis.
- **numpy**: For numerical operations.
- **matplotlib**: For data visualization.
- **seaborn**: For advanced data visualization.
- **scikit-learn**: For machine learning models and utilities like train-test split and vectorization.
- **re**: For text preprocessing using regular expressions.

## Data Source

The dataset used for this project is a CSV file named spam.csv that contains text messages labeled as "ham" (non-spam) or "spam." It is assumed that the file is hosted or uploaded within the working directory.

The dataset has the following columns:

1. **v1**: The category (ham/spam).
2. **v2**: The message content.

The CSV file is read and processed to prepare it for model training.

## 1. Importing Libraries

```
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns
```

These libraries are essential for data handling, numerical operations, visualization, and machine learning tasks.

## 2. Reading the Data

```
data = pd.read_csv('/content/spam.csv', encoding='ISO-8859-1')
```

The dataset is read into a pandas DataFrame. The encoding 'ISO-8859-1' is used to handle potential character encoding issues.

## 3. Data Inspection

```
print(data.head())
```

Displays the first few rows of the dataset to understand its structure and check the data quality.

## 4. Data Cleaning

- **Removing Unnecessary Columns:**

```
data = data.drop(columns=['Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'])
```

Columns that are unnecessary or empty (such as Unnamed: 2, Unnamed: 3, Unnamed: 4) are dropped from the dataset.

- **Renaming Columns:**

```
data = data.rename(columns={'v1': 'Category', 'v2': 'Message'})
```

The columns are renamed for better clarity: v1 becomes Category and v2 becomes Message.

## 5. Exploring Category Distribution

```
print(data['Category'].value_counts())
```

The Category column (ham/spam) is analyzed to check the distribution of labels.

## 6. Mapping Labels to Numeric Values

```
data['Category'] = data['Category'].map({'ham': 0, 'spam': 1})
```

The labels "ham" and "spam" are converted to numeric values (0 for ham and 1 for spam), as machine learning models work better with numerical data.

## 7. Handling Missing Data

```
print(data.isnull().sum())
```

The code checks for missing values in the dataset.

## 8. Removing Duplicates

```
data = data.drop_duplicates()
```

Duplicate rows are removed to ensure that the model trains on unique data.

## 9. Text Preprocessing

```
import re
```

```
data['Message'] = data['Message'].apply(lambda x: re.sub(r'\W+', ' ',  
x.lower()))
```


A regular expression (re.sub) is used to remove all non-word characters (e.g., punctuation) and convert the text to lowercase. To standardize the text data.

## 10. Exploratory Data Analysis (EDA)

- **Distribution of Spam vs. Ham:**

```
label_counts = data['Category'].value_counts()
plt.bar(['Ham', 'Spam'], label_counts, color=['lightblue', 'grey'])
plt.title('Spam vs Ham Distribution')
plt.xlabel('Label (Ham vs Spam)')
plt.ylabel('Count')
plt.show()
```

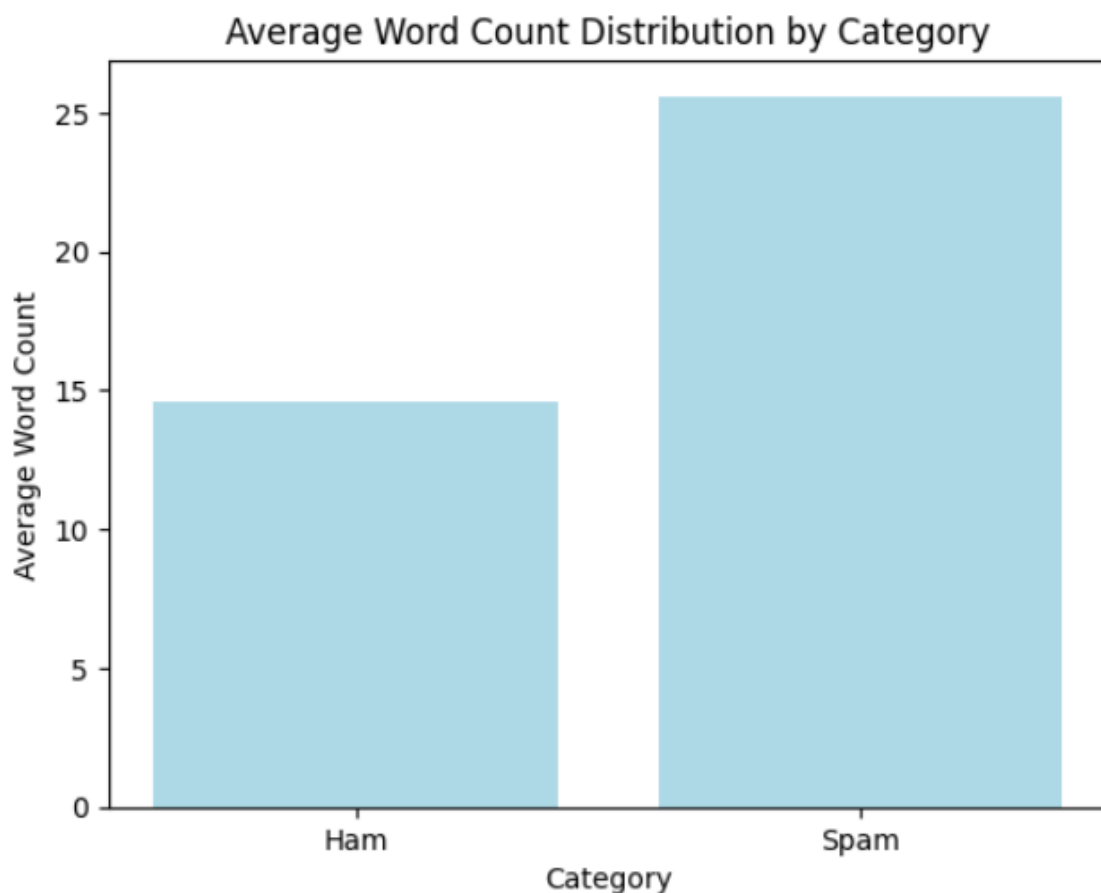
A bar plot is created to visualize the distribution of spam and ham messages in the dataset.



- **Average Word Count by Category:**

```
data['word_count'] = data['Message'].apply(lambda x: len(x.split()))  
word_count_by_category =  
data.groupby('Category')['word_count'].mean()  
plt.bar(['Ham', 'Spam'], word_count_by_category.values,  
color='lightblue')  
plt.xlabel('Category')  
plt.ylabel('Average Word Count')  
plt.title('Average Word Count Distribution by Category')  
plt.show()
```

The average word count of messages in each category (spam and ham) is plotted to observe if there is a significant difference.



## 11. Text Vectorization

```
from sklearn.feature_extraction.text import TfidfVectorizer  
tfidf = TfidfVectorizer(max_features=5000)  
X = tfidf.fit_transform(X).toarray()
```

The text data is transformed into numerical data using the **TF-IDF** (Term Frequency-Inverse Document Frequency) vectorizer, limiting to the 5000 most important words.

## 12. Data Splitting and Train-Test Split

```
X = data['Message']  
y = data['Category']  
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,  
random_state=42)
```

The dataset is split into training and testing sets, with 20% of the data used for testing.

## 13. Model Training

```
from sklearn.naive_bayes import MultinomialNB  
model = MultinomialNB()  
model.fit(X_train, y_train)
```

The *Multinomial Naive Bayes* classifier is trained on the training data.

## 14. Model Evaluation

```
from sklearn.metrics import accuracy_score, confusion_matrix  
print("Accuracy:", accuracy_score(y_test, y_pred))  
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

The accuracy of the model is evaluated, and a confusion matrix is generated to show the number of correct and incorrect predictions.

- **Confusion Matrix Visualization:**

```
cm = confusion_matrix(y_test, y_pred)  
sns.heatmap(cm, annot=True, fmt='d', cmap='coolwarm')  
plt.title('Confusion Matrix')  
plt.xlabel('Predicted')  
plt.ylabel('Actual')  
plt.show()
```

A heatmap is created to visually represent the confusion matrix.

## **Conclusion**

- The model was trained and tested on a spam message dataset.
- The Multinomial Naive Bayes classifier performed well, as shown by the accuracy and confusion matrix.
- Future improvements could involve experimenting with different models, hyperparameter tuning, or incorporating more complex text preprocessing techniques.